

Introduction to PostgreSQL

Brief Introduction to PostgreSQL

PostgreSQL is an advanced open-source object-relational database management system (ORDBMS) known for its robustness, extensibility, and adherence to SQL standards. It offers a wide range of features and capabilities that make it a popular choice for both small-scale applications and large-scale enterprise systems.

Key Features of PostgreSQL

- 1. Relational Database Management:** PostgreSQL is a powerful relational database management system that allows you to store, retrieve, and manipulate structured data efficiently.
- 2. ACID Compliance:** PostgreSQL ensures ACID (Atomicity, Consistency, Isolation, Durability) compliance, which guarantees data integrity and reliability even in the presence of concurrent transactions.
- 3. Data Types:** PostgreSQL provides an extensive range of built-in and user-defined data types, including numeric, string, boolean, date/time, JSON, XML, and more. This flexibility enables you to model and store complex data structures.
- 4. Extensibility:** PostgreSQL allows you to extend its functionality through user-defined functions, data types, operators, and even procedural languages. This extensibility enables customisations and the creation of advanced database solutions.
- 5. Concurrency Control:** PostgreSQL utilises Multi-Version Concurrency Control (MVCC) to manage concurrent access to data, ensuring consistent and isolated transactions while allowing high concurrency.
- 6. Advanced Indexing:** PostgreSQL offers various indexing techniques, including B-tree, hash, GiST (Generalised Search Tree), GIN (Generalised Inverted Index), and SP-GiST. These indexes enhance query performance and enable efficient data retrieval.
- 7. Rich Querying Capabilities:** PostgreSQL supports complex SQL queries, subqueries, joins, window functions, common table expressions (CTEs), full-text search, and geospatial data operations. It also provides stored procedures, triggers, and user-defined functions to build powerful database logic.
- 8. Replication and High Availability:** PostgreSQL supports various replication methods, including streaming replication, logical replication, and synchronous replication. These features ensure data redundancy and high availability in distributed database setups.
- 9. Foreign Data Wrappers:** PostgreSQL allows you to connect and interact with data from external data sources using foreign data wrappers (FDWs). It enables seamless integration with other databases or systems.
- 10. Security and Authentication:** PostgreSQL provides robust security mechanisms, including role-based access control (RBAC), SSL encryption, password authentication, and support for LDAP, Kerberos, and other authentication methods.

Community and Ecosystem

PostgreSQL has a vibrant and active community of developers and users who contribute to its development, documentation, and support. The PostgreSQL Global Development Group oversees the project's development and ensures its ongoing improvement and stability.

Additionally, PostgreSQL has a rich ecosystem of tools and extensions that extend its functionality. These include graphical administration tools (such as pgAdmin and DBeaver), data modeling tools, connection libraries for various programming languages, and numerous extensions for specialised use cases.

PostgreSQL's open-source nature and its compatibility with various operating systems (including Linux, Windows, and macOS) make it a versatile and widely adopted database management system across different industries and applications.

Installation and Setup

Downloading and installing PostgreSQL

1. Visit the official PostgreSQL website (<https://www.postgresql.org>) and navigate to the "Download" section.
2. Choose the appropriate version for your operating system and download the installer.
3. Run the installer and follow the on-screen instructions to complete the installation process.
4. During the installation, you can choose the installation directory, port number, and other configuration options based on your preferences.

Configuring PostgreSQL

1. After the installation, locate the PostgreSQL configuration files. The main configuration file is typically named `postgresql.conf`.
2. Open the configuration file in a text editor and modify the settings as needed. Some common configuration options include adjusting the memory allocation, specifying the data directory, and enabling or disabling certain features.
3. Save the changes to the configuration file.

Connecting to PostgreSQL

1. Start the PostgreSQL service or server. The method depends on your operating system. For example, on Linux, you can use the following command:

```
1 sudo service postgresql start
```

2. To connect to the PostgreSQL database, you can use the `psql` command-line tool or a graphical client such as pgAdmin.
3. For the `psql` command-line tool, open a terminal or command prompt and run the following command:

```
1 psql -U <username> -d <database_name> -h <host> -p <port>
```

Replace `<username>`, `<database_name>`, `<host>`, and `<port>` with the appropriate values. Enter the password when prompted.

4. If the connection is successful, you will see the PostgreSQL prompt (`psql`).

PostgreSQL Fundamentals

Databases and Tables

Creating a database

1. To create a new database, you can use the `CREATE DATABASE` statement. For example:

```
1 CREATE DATABASE dbname;
```

Replace `dbname` with the desired name of your database.

2. You can also specify additional options while creating the database, such as the owner, character set, and collation.

```
1 CREATE DATABASE dbname
2 OWNER username
3 ENCODING 'UTF8'
```

```
4 LC_COLLATE = 'en_US.UTF-8'  
5 LC_CTYPE = 'en_US.UTF-8';
```

Switching between databases

1. To switch to a different database within the `psql` command-line tool, use the following command:

```
1 \c dbname
```

Replace `dbname` with the name of the database you want to connect to.

Creating tables and defining columns

1. To create a new table, you can use the `CREATE TABLE` statement. Specify the table name and column definitions. For example:

```
1 CREATE TABLE tablename (  
2     column1 datatype,  
3     column2 datatype,  
4     ...  
5 );
```

Replace `tablename` with the desired name for your table. Specify the column names and their respective data types.

2. You can also add constraints, such as primary keys, unique constraints, and foreign keys, to the table definition.

Data types in PostgreSQL

1. To modify an existing table, you can use the `ALTER TABLE` statement. It allows you to add, modify, or drop columns, constraints, and other table properties.
2. For example, to add a new column to an existing table:

```
1 ALTER TABLE tablename  
2 ADD COLUMN newcolumn datatype;
```

Replace `tablename` with the name of the table, and specify the new column name and data type.

Dropping Tables

1. To drop a table, use the `DROP TABLE` statement. This permanently deletes the table and all its data. Exercise caution when using this command as it cannot be undone.

```
1 DROP TABLE tablename;
```

Replace `tablename` with the name of the table you want to drop.

2. If the table has dependencies, such as foreign keys or views, you may need to use the `CASCADE` option to drop those dependencies along with the table.

```
1 DROP TABLE tablename CASCADE;
```

Viewing Tables

1. To view the list of tables in the current database, you can query the `pg_tables` catalog table.

```
1 SELECT tablename FROM pg_tables WHERE schemaname = 'public';
```

This query retrieves all table names from the `public` schema. Adjust the schema name as per your requirements.

2. You can also use the `\dt` command in the `psql` command-line tool to list all tables in the current database.

Modifying Table Data

1. To insert data into a table, use the `INSERT INTO` statement.

```
1 INSERT INTO tablename (column1, column2, ...)
2 VALUES (value1, value2, ...);
```

2. To update existing data in a table, use the `UPDATE` statement.

```
1 UPDATE tablename
2 SET column1 = newvalue1, column2 = newvalue2, ...
3 WHERE condition;
```

3. To delete data from a table, use the `DELETE FROM` statement.

```
1 DELETE FROM tablename
2 WHERE condition;
```

Querying Tables

1. To retrieve data from a table, use the `SELECT` statement.

```
1 SELECT column1, column2, ...
2 FROM tablename
3 WHERE condition;
```

2. You can use various clauses with the `SELECT` statement, such as `WHERE` for filtering, `ORDER BY` for sorting, `GROUP BY` for grouping, and `JOIN` for combining data from multiple tables.

Adding constraints to tables

Constraints in PostgreSQL are used to enforce rules and maintain data integrity within tables. They define limits and conditions on the values that can be inserted, updated, or deleted from specific columns. Here are two common types of constraints and examples of how to add them:

1. A primary key constraint ensures the uniqueness and non-nullability of a column or a combination of columns. It uniquely identifies each row in a table.

```
1 -- Example: Adding a primary key constraint to a column
2 ALTER TABLE tablename
3 ADD PRIMARY KEY (columnname);
```

2. A foreign key constraint establishes a relationship between two tables by enforcing referential integrity. It ensures that values in a column (or a combination of columns) of one table match the values in the referenced column(s) of another table.

```
1 -- Example: Adding a foreign key constraint to a column
2 ALTER TABLE childtable
3 ADD CONSTRAINT fk_constraint_name
4 FOREIGN KEY (columnname)
5 REFERENCES parenttable (referencedcolumn);
```

In the above example, `childtable` is the table where the foreign key is being added, `fk_constraint_name` is the name of the foreign key constraint, `columnname` is the column in `childtable` referencing the `referencedcolumn` in `parenttable`.

Indexes

Indexes in PostgreSQL are used to improve query performance by providing efficient access to data. They are data structures that allow quick lookups, sorting, and filtering of data based on the indexed columns. PostgreSQL supports various types of indexes, including B-tree indexes, hash indexes, GiST indexes, GIN indexes, and SP-GiST indexes.

In this section, we'll focus on Primary, Composite, Partial indexes. Then we'll look at B-tree indexes which are the most commonly used index type in PostgreSQL.

Primary key indexes

In PostgreSQL, a primary key is a constraint that uniquely identifies each row in a table. It ensures that the values in the specified column or columns are unique and cannot be null. A primary key index is automatically created when you define a primary key on a table. It provides fast and efficient access to data based on the primary key column(s).

Creating a Primary Key Index

To create a primary key index in PostgreSQL, you can use the `PRIMARY KEY` constraint while defining the table. Here's an example:

```
1 CREATE TABLE tablename (  
2     column1 datatype,  
3     column2 datatype,  
4     ...  
5     PRIMARY KEY (column1)  
6 );
```

In the above example, `tablename` is the name of the table, `column1` is the column that serves as the primary key, and `datatype` represents the appropriate data type for the column.

Using a Primary Key Index

Once a primary key index is created, you can utilise it in various ways:

1. **Uniqueness enforcement:** The primary key index ensures the uniqueness of values in the specified column(s). If you attempt to insert or update a row with a duplicate primary key value, PostgreSQL will raise an error.
2. **Fast data retrieval:** The primary key index provides efficient access to data based on the primary key column(s). When you execute a query that involves the primary key column(s), PostgreSQL can utilise the index to quickly locate the corresponding row(s).
3. **Implicit index usage:** When you perform operations that require the primary key, such as `UPDATE`, `DELETE`, or `MERGE` statements, PostgreSQL automatically utilises the primary key index to locate the affected row(s) efficiently.
4. **Referential integrity enforcement:** Primary key indexes can be referenced by foreign key constraints in other tables. This ensures that foreign key references to the primary key values are valid and consistent.

Remember to choose primary key column(s) that are meaningful and unique to your data model. It's important to use primary key indexes effectively to maintain data integrity and optimise query performance.

Composite indexes

In PostgreSQL, a composite index is an index that is created on multiple columns of a table. It allows you to optimize queries that involve conditions or joins on multiple columns. Composite indexes can improve query performance by reducing the number of disk accesses needed to satisfy the query.

Creating a Composite Index

To create a composite index, you specify multiple columns within the index definition. Here's an example:

```
1 CREATE INDEX indexname ON tablename (column1, column2, ...);
```

In the above example, `indexname` is the name of the index, `tablename` is the name of the table, and `column1`, `column2`, etc., are the columns included in the composite index.

Benefits of Composite Indexes

1. **Improved Query Performance:** Composite indexes can significantly improve query performance for queries that involve conditions or joins on multiple columns. By having an index on multiple columns used in a query, PostgreSQL can quickly locate the relevant rows without performing a full table scan.
2. **Index Selectivity:** Composite indexes can provide higher selectivity than individual indexes on single columns. They can capture more specific combinations of column values, allowing the optimizer to make better choices when executing queries.
3. **Index-Only Scans:** In certain cases, a composite index can cover all the columns required by a query. This enables PostgreSQL to perform an index-only scan, accessing the data directly from the index without having to access the table, resulting in improved query performance.

Considerations for Composite Indexes

When working with composite indexes, it's important to keep the following considerations in mind:

1. **Column Order:** The order of columns in a composite index is significant. The most selective column(s) should be placed first in the index definition. This helps maximize the effectiveness of the index for a wider range of queries.
2. **Query Patterns:** Analyze the query patterns and their corresponding column combinations. If a specific order of columns is consistently used in queries, create a composite index with those columns in the same order. This can optimize query performance.
3. **Column Cardinality:** Consider the cardinality (distinct values) of the columns in the composite index. Higher cardinality columns generally provide more selective indexes.
4. **Index Maintenance:** Composite indexes require additional storage space and maintenance overhead compared to individual indexes. Be mindful of the impact on insert, update, and delete operations when deciding to create composite indexes.

Remember to carefully analyze your query patterns and consider the trade-offs between index size, maintenance overhead, and query performance when deciding to create composite indexes.

Take a look at some of the best practices for creating composite indexes: [Best practices for creating multi-column indexes in PostgreSQL](#)



Partial indexes

In PostgreSQL, a partial index is an index that is created on a subset of rows in a table, based on a specified condition. It allows you to create an index on a specific subset of data that is relevant to your queries, thereby improving query performance and reducing index size.

Creating a Partial Index

To create a partial index, you specify a condition using the `WHERE` clause while defining the index. Here's an example:

```
1 CREATE INDEX indexname ON tablename (column)
2 WHERE condition;
```

In the above example, `indexname` is the name of the index, `tablename` is the name of the table, `column` is the column on which the index is created, and `condition` is the condition that determines which rows are included in the index.

Benefits of Partial Indexes

1. **Improved Query Performance:** Partial indexes allow you to create indexes on a subset of data that is relevant to your queries. By indexing only the necessary rows, you can significantly improve query performance for specific queries that match the index condition.
2. **Reduced Index Size:** Since partial indexes only include a subset of rows, they occupy less disk space compared to indexes on the entire table. This can be especially beneficial when dealing with large tables and selective queries.
3. **Faster Index Maintenance:** Partial indexes require less maintenance during data modifications (inserts, updates, and deletes) compared to indexes on the entire table. This can result in faster data modifications and reduced overhead.

Considerations for Partial Indexes

When working with partial indexes, consider the following:

1. **Selectivity of Condition:** The condition used to create the partial index should be highly selective. It should match a relatively small subset of rows to ensure that the index is effective in improving query performance.
2. **Query Patterns:** Analyse the query patterns and identify queries that can benefit from a partial index. Create the partial index based on the conditions used in those queries to optimize their performance.
3. **Index Maintenance:** Partial indexes require additional maintenance overhead during data modifications on the indexed subset of rows. Consider the impact on insert, update, and delete operations when deciding to create partial indexes.
4. **Index Usage:** PostgreSQL's query planner will automatically consider using a partial index when it is applicable to a query. However, it's important to monitor and analyze query plans to ensure that the partial index is being utilized effectively.

Partial indexes are a powerful feature in PostgreSQL that allow you to optimize query performance and reduce index size for specific subsets of data. Carefully analyze your query patterns and consider the selectivity of the condition before creating partial indexes.

Now lets take a look at B-tree indexes. They are the most common type of indexes which we use on postgres.

B-tree Indexes

B-tree indexes in PostgreSQL are balanced tree structures that store sorted key-value pairs. They are efficient for most types of queries and can handle both equality and range searches efficiently. B-tree indexes are automatically created when a primary key or unique constraint is defined on a table, but they can also be manually created on any column(s) for improved query performance.

Key Concepts

1. **B-tree Structure:** B-trees are self-balancing tree structures that maintain sorted data. In a B-tree index, each level of the tree (except the leaf level) contains multiple key-value pairs, and the leaf level contains pointers to actual table rows.
2. **Index Key:** The indexed column(s) or expression used to create the index is known as the index key. It determines the sorting order of the index.
3. **Index Scan:** When a query involves a search condition that can be matched by the index key, PostgreSQL can use the B-tree index for an index scan. It traverses the index tree to find the matching entries efficiently.

Creating and Using B-tree Indexes

To create a B-tree index in PostgreSQL, you can use the `CREATE INDEX` statement. Here's an example:

```
1 CREATE INDEX idx_name ON table_name (column1, column2);
```

To use an index in a query, you typically include the indexed columns in the `WHERE` clause of your query. PostgreSQL's query planner can then choose to utilize the index for faster data retrieval.

Benefits and Considerations

B-tree indexes offer several benefits:

- **Improved query performance:** B-tree indexes allow faster data retrieval, especially for equality and range queries involving the indexed columns.
- **Sorted data:** The B-tree structure keeps the index keys in sorted order, enabling efficient sorting operations.

However, there are some considerations to keep in mind:

- **Index maintenance:** As data is inserted, updated, or deleted in the table, the corresponding index entries need to be modified, which incurs some overhead. Regular maintenance, such as running `VACUUM` or `autovacuum`, helps optimize index performance.
- **Index size:** B-tree indexes consume disk space to store the index structure and key-value pairs. Large indexes may have a significant disk footprint.

Additional Indexing Features

PostgreSQL's B-tree indexes offer additional features for optimising query performance and flexibility:

- **Partial Indexes:** You can create an index that covers only a subset of rows in a table based on a defined condition. This can help reduce index size and improve query performance for specific subsets of data.
- **Expression Indexes:** PostgreSQL allows creating indexes on expressions or function results, enabling faster lookups based on computed values.
- **Multi-column Indexes:** B-tree indexes can span multiple columns. Creating indexes on combinations of columns commonly used in queries can further enhance performance.
- **Index-Only Scans:** Under certain conditions, PostgreSQL can perform index-only scans, retrieving data directly from the index without accessing the table, resulting in faster query execution.

Query Optimisation and Performance

In PostgreSQL, `VACUUM` and `ANALYZE` are two essential commands used for maintaining table statistics and reclaiming disk space.

Regularly performing `VACUUM` and `ANALYZE` operations, either manually or through `autovacuum`, helps to maintain optimal performance, prevent bloat, and ensure accurate query planning in PostgreSQL.

VACUUM

The `VACUUM` command is used to reclaim disk space and optimize table storage in PostgreSQL. It performs several tasks, including:

1. **Freeing up space:** When data is updated or deleted, PostgreSQL marks the old versions of rows as dead but does not immediately reclaim the disk space they occupy. `VACUUM` removes these dead rows and frees up space, ensuring efficient storage utilization.
2. **Updating visibility information:** PostgreSQL uses a multi-version concurrency control (MVCC) mechanism to manage concurrent transactions. `VACUUM` updates the visibility information, allowing PostgreSQL to determine which rows are visible to different transactions correctly.
3. **Preventing transaction ID wraparound:** PostgreSQL uses a transaction ID (XID) counter to track the age of transactions. If the XID counter approaches its maximum value, a transaction ID wraparound can occur, leading to database corruption. `VACUUM` prevents this by updating the transaction log and recycling old transaction IDs.

The `VACUUM` command can be executed at the database, schema, or table level. Here's an example of running `VACUUM` on a table:


```
1 VACUUM tablename;
```

ANALYZE

The ANALYZE command is used to collect statistics about tables and indexes in PostgreSQL. It helps the query optimizer make informed decisions when generating query plans, leading to better performance. The main tasks performed by ANALYZE include:

1. **Collecting table and column statistics:** ANALYZE examines the data distribution within tables and columns, collecting information such as the number of distinct values, data correlation, and distribution histograms. These statistics enable the query optimizer to estimate the selectivity of conditions and choose the most efficient execution plans.
2. **Updating system catalogs:** ANALYZE updates the system catalogs with the collected statistics, ensuring that the query optimizer has the most up-to-date information for query planning.

The ANALYZE command can be executed at the database, schema, or table level. Here's an example of running ANALYZE on a table:

```
1 ANALYZE tablename;
```

AUTOVACUUM

PostgreSQL also provides an AutoVacuum mechanism that automatically performs VACUUM and ANALYZE operations on tables and indexes. AutoVacuum continuously monitors the system and triggers VACUUM and ANALYZE processes based on predefined thresholds and configuration settings.

By default, AutoVacuum is enabled in PostgreSQL, and it is recommended to keep it enabled to ensure the regular maintenance of tables and indexes. However, it's essential to monitor and tune the AutoVacuum settings to suit the specific needs and workload characteristics of your database.

Here are a few PostgreSQL catalog tables and functions that you can use to check and monitor AutoVacuum.

1. pg_stat_all_tables

The `pg_stat_all_tables` catalog view provides information about autovacuum activities and statistics for all tables in the current database. You can query this view to obtain details such as the number of live and dead rows, the last autovacuum time, and the number of times autovacuum has been triggered for each table.

```
1 SELECT * FROM pg_stat_all_tables;
```

2. pg_stat_user_tables

Similar to `pg_stat_all_tables`, the `pg_stat_user_tables` catalog view provides autovacuum statistics for tables, but only for tables owned by the current user. This view can be used to monitor autovacuum activity specifically for the tables you own.

```
1 SELECT * FROM pg_stat_user_tables;
```

3. pg_stat_progress_vacuum

The `pg_stat_progress_vacuum` catalog view provides real-time information about ongoing vacuum operations, including autovacuum. It shows details such as the current phase of the vacuum process, the number of scanned and vacuumed pages, and the current table being vacuumed.

```
1 SELECT * FROM pg_stat_progress_vacuum;
```

EXPLAIN

In PostgreSQL, the EXPLAIN command is used to analyse the execution plan of a query. It provides valuable insights into how the PostgreSQL query optimiser plans to execute the query and helps identify potential performance bottlenecks.

Basic Usage

To use the EXPLAIN command, you simply prefix your query with EXPLAIN. For example:

```
1 EXPLAIN SELECT column1, column2 FROM tablename WHERE condition;
```

This command will output the query plan, which includes information about the order of operations, the join methods, and the access methods chosen by the query planner.

Query Plan Output

The output of the EXPLAIN command provides a detailed breakdown of the query plan. Each row in the output represents an operation performed during query execution, known as a "plan node." The important columns in the output include:

- **Node Type:** The type of operation being performed, such as "Seq Scan" (sequential scan), "Index Scan" (index scan), "Hash Join," etc.
- **Join Type:** For join operations, this column indicates the type of join being used, such as "Nested Loop," "Hash," or "Merge."
- **Relation Name:** The name of the table or index being accessed.
- **Filter/Join Conditions:** The conditions used for filtering or joining the data.
- **Cost:** The estimated cost of executing the operation, which is used by the query planner to choose the most efficient plan.
- **Actual Rows:** The number of rows processed or returned by each plan node during execution.

By carefully analysing the output of EXPLAIN, you can identify potential performance issues, such as sequential scans instead of index scans, high-cost operations, or incorrect join strategies. This allows you to optimise the query by adding appropriate indexes, rewriting the query, or adjusting the configuration.

Additional Options

The EXPLAIN command provides additional options to obtain more detailed information:

- **EXPLAIN ANALYZE:** This option not only outputs the query plan but also executes the query and provides actual run-time statistics, such as the actual execution time, number of rows returned, and memory usage. It gives you a more accurate representation of query performance.
- **EXPLAIN VERBOSE:** This option provides more detailed information about each plan node, including additional properties, statistics, and condition details.

EXPLAIN ANALYZE

In PostgreSQL, the EXPLAIN ANALYZE command combines the functionality of the EXPLAIN command with the actual execution of the query. It provides detailed query plan information along with runtime statistics, allowing you to analyze both the expected plan and the actual performance of the query.

Basic Usage

To use the EXPLAIN ANALYZE command, you prefix your query with EXPLAIN ANALYZE instead of just EXPLAIN. For example:

```
1 EXPLAIN ANALYZE SELECT column1, column2 FROM tablename WHERE condition;
```

When you execute this command, PostgreSQL not only provides the query plan but also executes the query and gathers runtime statistics.

Query Plan Output with Runtime Statistics

The output of EXPLAIN ANALYZE includes the same information as the EXPLAIN command, such as the plan nodes, join types, relations, costs, and filter/join conditions. Additionally, it provides the following runtime statistics:

- **Actual Execution Time:** The actual time taken to execute the query, including the time spent on planning, data retrieval, and other operations.
- **Actual Rows:** The actual number of rows processed or returned by each plan node during execution.
- **Actual Loops:** The number of repetitions performed by a looped operation, such as nested loops or joins.
- **Actual Total Time:** The total time spent executing a plan node, including all repetitions or loops.
- **Peak Memory Usage:** The maximum amount of memory used by the query execution.

These runtime statistics give you a more accurate understanding of the query performance based on the actual execution. They can help identify performance bottlenecks, such as slow operations, unexpected row counts, or excessive memory consumption.

Interpreting the Output

When using EXPLAIN ANALYZE, you should pay attention to the actual execution time, row counts, and other statistics. Compare them with the estimated values from the EXPLAIN command to identify discrepancies and potential areas for optimization.

By analyzing the runtime statistics, you can gain insights into the efficiency of your query plan, identify performance issues, and make informed decisions to improve query performance.

Caution

Keep in mind that running EXPLAIN ANALYZE on large or resource-intensive queries can have an impact on the database server's performance. It is advisable to use EXPLAIN ANALYZE judiciously, especially in production environments, and consider running it on representative subsets of data if possible.

The EXPLAIN ANALYZE command is a powerful tool for query optimization and performance tuning in PostgreSQL. It allows you to assess the actual runtime behaviour of your queries, helping you make informed decisions to optimize their performance.

Monitoring and Statistics

The pg_catalog Schema: Accessing System Information

The PostgreSQL system catalog, also known as the pg_catalog schema, contains a collection of system tables and views that store metadata about the database objects and system information. Here are some basic things in the pg_catalog:

1. **pg_tables:** This catalog table contains information about all tables in the current database, such as table names, schemas, and owners.
2. **pg_views:** It stores information about views in the database, including view names, definitions, and owners.
3. **pg_indexes:** This catalog table contains information about indexes defined on tables, including the index name, table name, and index definition.
4. **pg_columns:** It stores details about the columns of all tables in the database, including the column name, data type, length, and whether it allows null values.
5. **pg_constraint:** This catalog table contains information about constraints defined on tables, such as primary keys, unique constraints, foreign keys, and check constraints.
6. **pg_stat_all_tables:** It provides statistics about all tables in the database, including the number of rows, disk size, and other useful metrics for performance analysis.

7. **pg_stat_user_tables**: This view is similar to `pg_stat_all_tables` but provides statistics only for tables owned by the current user.
8. **pg_stat_all_indexes**: It provides statistics about the indexes in the database, including the number of index scans, index size, and index hit rates.
9. **pg_stat_user_indexes**: Similar to `pg_stat_all_indexes`, this view provides statistics only for indexes owned by the current user.
10. **pg_stat_activity**: This view contains information about the current activity of database sessions, including the client address, query being executed, and the current state of the session.
11. **pg_namespace**: It stores information about database schemas, including schema names and their owners.
12. **pg_settings**: Contains configuration settings and their current values for the PostgreSQL server.
13. **pg_class**: Stores information about tables, views, and indexes, such as their names, sizes, and access methods.
14. **pg_namespace**: Contains details about database schemas, including schema names and owner information.
15. **pg_attribute**: Stores information about table columns, including their names, types, constraints, and other attributes.
16. **pg_stat_replication**: If you have configured replication in PostgreSQL, the `pg_stat_replication` view provides information about the replication connections. It includes details such as the status of replication, replication lag, and other statistics related to replication processes.
17. **pg_foreign_server**: This catalog table stores information about foreign servers defined in PostgreSQL. Foreign servers are used for establishing connections to remote databases or systems. You can query this table to check the configured foreign servers and their connection details.
18. **pg_foreign_table**: If you have foreign tables defined in your database, the `pg_foreign_table` catalog table contains information about those tables. It includes details such as the foreign server associated with the table and the options used for the foreign table connection.

These are just a few examples of the tables and views available in the `pg_catalog` schema. Exploring the `pg_catalog` allows you to gather valuable insights about the structure, statistics, and activity of your PostgreSQL database.

Query Examples

Here are a few examples of querying the `pg_catalog` schema to retrieve system information:

1. Get a list of all tables in the current database:

```
1 SELECT tablename FROM pg_catalog.pg_tables WHERE schemaname = 'public';
```

2. Retrieve the column names and types of a specific table:

```
1 SELECT column_name, data_type FROM pg_catalog.pg_columns WHERE table_name = 'tablename';
```

3. Retrieve the index information for a table:

```
1 SELECT indexname, indexdef FROM pg_catalog.pg_indexes WHERE tablename = 'tablename';
```

4. Retrieve the current values of configuration settings:

```
1 SELECT name, setting FROM pg_catalog.pg_settings;
```

These are just a few examples of how you can query the `pg_catalog` schema to access system information. The catalog tables and views in the `pg_catalog` schema provide a wealth of information about the database system, which can be useful for various administrative tasks, performance monitoring, and troubleshooting purposes.

pg_stat_activity: Monitoring Database Activity

The `pg_stat_activity` view is a system catalog view in PostgreSQL that provides information about the currently active connections and activities in the database server. It allows you to monitor and analyze the ongoing activity in the database, including executing queries, waiting for locks, and performing administrative tasks.

Querying pg_stat_activity

You can query the `pg_stat_activity` view to retrieve valuable information about the active connections and their associated activities. Some of the important columns in this view include:

- **datid**: The OID (object identifier) of the database where the activity is taking place.
- **datname**: The name of the database where the activity is taking place.
- **pid**: The process ID of the backend server process handling the connection.
- **username**: The name of the user who initiated the connection.
- **client_addr**: The IP address of the client machine.
- **client_port**: The port number of the client connection.
- **backend_start**: The timestamp when the backend server process started.
- **query_start**: The timestamp when the currently executing query or command started.
- **state**: The current state of the connection, such as "active", "idle", or "idle in transaction".
- **query**: The currently executing query or command (truncated to a certain length).
- **wait_event**: The event for which the connection is waiting, if applicable.
- **wait_event_type**: The type of event for which the connection is waiting, if applicable.

By querying the `pg_stat_activity` view, you can gain insights into the current workload, identify long-running or blocked queries, monitor connections, and track resource usage.

Example Queries

Here are a few examples of querying the `pg_stat_activity` view to monitor database activity:

1. Retrieve the currently running queries and their associated details:

```
1 SELECT pid, query, state, query_start FROM pg_stat_activity WHERE state = 'active';
```

2. Find the connections waiting on a specific event:

```
1 SELECT pid, wait_event, wait_event_type FROM pg_stat_activity WHERE wait_event IS NOT NULL;
```

3. Get a list of idle connections:

```
1 SELECT pid, username, state, backend_start FROM pg_stat_activity WHERE state = 'idle';
```

These are just a few examples of how you can utilize the `pg_stat_activity` view to monitor database activity. By regularly querying this view, you can gain insights into the current state of the database, track query execution, identify performance bottlenecks, and troubleshoot issues.

Common functions used for monitoring in Postgres

1. pg_stat_get_activity(queryid)

- This function returns information about the activity of a specific backend process identified by the `queryid`.
- Example usage:

```
1 SELECT * FROM pg_stat_get_activity(12345);
```

2. `pg_stat_get_backend_pid()`

- This function returns the process ID (PID) of the current backend process.
- Example usage:

```
1 SELECT pg_stat_get_backend_pid();
```

3. `pg_stat_reset()`

- This function resets all statistics collected by the PostgreSQL statistics collector.
- Example usage:

```
1 SELECT pg_stat_reset();
```

4. `pg_stat_clear_snapshot()`

- This function clears the current snapshot used for collecting statistics, allowing new statistics to be collected.
- Example usage:

```
1 SELECT pg_stat_clear_snapshot();
```

5. `pg_stat_get_db_numbackends(dbid)`

- This function returns the number of currently active backends connected to a specific database identified by the `dbid`.
- Example usage:

```
1 SELECT pg_stat_get_db_numbackends(12345);
```

6. `pg_stat_get_db_xact_commit(dbid)`

- This function returns the number of transactions committed in a specific database identified by the `dbid`.
- Example usage:

```
1 SELECT pg_stat_get_db_xact_commit(12345);
```

7. `pg_stat_get_db_size(dbid)`

- This function returns the size of a specific database identified by the `dbid`.
- Example usage:

```
1 SELECT pg_stat_get_db_size(12345);
```

8. `pg_stat_file(filename)`

- This function returns statistical information about a specific file in the PostgreSQL data directory. It provides details such as the number of blocks, size, and modification timestamp.
- Example usage:

```
1 SELECT * FROM pg_stat_file('path/to/file');
```

9. `pg_stat_get_live_tuples(relid)`

- This function returns the number of live tuples (rows) in a specific table identified by the `relid`.
- Example usage:

```
1 SELECT pg_stat_get_live_tuples('tablename'::regclass);
```

10. `pg_stat_get_dead_tuples(relid)`

- This function returns the number of dead tuples (rows) in a specific table identified by the `relid`.
 - Example usage:
-

```
1 SELECT pg_stat_get_dead_tuples('tablename'::regclass);
```

11. `pg_stat_get_blocks_fetched(relid)`

- This function returns the number of disk blocks fetched from a specific table identified by the `relid`.
- Example usage:

```
1 SELECT pg_stat_get_blocks_fetched('tablename'::regclass);
```

Internal Workings

Multi-Version Concurrency Control (MVCC)

MVCC is a technique used in PostgreSQL to manage concurrent transactions and provide high concurrency in a multi-user database environment. It ensures that each transaction sees a consistent snapshot of the database, even when multiple transactions are accessing or modifying the same data simultaneously.

Overview of MVCC in PostgreSQL

In MVCC, instead of locking entire tables or rows, PostgreSQL creates multiple versions of a row, each representing a different state of the data at a particular point in time. These versions are maintained in a way that allows concurrent transactions to read and modify data without conflicting with each other.

Each transaction in PostgreSQL has its own snapshot of the database, which includes a set of database versions as they existed at the start of the transaction. When a transaction reads data, it retrieves the appropriate version of the data based on its snapshot. This ensures that the transaction sees a consistent view of the database, even if other transactions are modifying the data concurrently.

How MVCC Manages Concurrent Transactions

MVCC in PostgreSQL utilizes the following mechanisms to manage concurrent transactions:

1. **Versioning:** When a row is updated or deleted, a new version of the row is created, while the old version is retained. Each version is associated with a transaction ID (XID) to identify the creating transaction.
2. **Transaction ID (XID):** Every transaction in PostgreSQL is assigned a unique transaction ID. This allows PostgreSQL to determine which versions of the data should be visible to each transaction based on its snapshot.
3. **Visibility Rules:** The visibility of data versions is determined by comparing the transaction's snapshot with the XIDs associated with each data version. A version is visible to a transaction if its XID is older than the transaction's snapshot XID. This ensures that a transaction only sees data that was committed before its start.

Transaction Isolation Levels in PostgreSQL

MVCC provides read consistency and supports different transaction isolation levels in PostgreSQL. The transaction isolation levels determine the level of isolation and concurrency guarantees provided by PostgreSQL.

PostgreSQL supports multiple transaction isolation levels that determine the level of concurrency and data consistency provided by the database. Each isolation level offers a different trade-off between concurrency and data integrity, allowing you to choose the appropriate

level based on your application's requirements.

The supported transaction isolation levels in PostgreSQL are:

1. Read Uncommitted

The Read Uncommitted isolation level provides the lowest level of isolation and allows transactions to read uncommitted changes made by other concurrent transactions. This means that a transaction can see dirty, uncommitted data. Read Uncommitted offers high concurrency but sacrifices data consistency.

2. Read Committed

The Read Committed isolation level ensures that a transaction sees only committed data. It guarantees that any data read by a transaction is committed and provides a higher level of consistency compared to Read Uncommitted. However, it doesn't guarantee consistency between multiple reads within the same transaction.

3. Repeatable Read

The Repeatable Read isolation level guarantees that within a transaction, the same query always returns the same data, regardless of concurrent changes made by other transactions. It ensures that any data read by a transaction remains consistent throughout the transaction's duration. However, it allows non-repeatable reads and phantom reads, meaning that new rows may be inserted or existing rows may be deleted by other concurrent transactions.

4. Serializable

The Serializable isolation level provides the highest level of isolation and guarantees that concurrent transactions have the same effect as if they were executed serially, in a predefined order. It eliminates the possibility of dirty reads, non-repeatable reads, and phantom reads by acquiring locks to prevent concurrent modifications. Serializable provides strong data consistency but can potentially result in serialization failures, requiring transactions to be retried.

Choosing the Right Isolation Level

Choosing the appropriate transaction isolation level depends on the requirements of your application. Consider the following factors:

- **Concurrency:** Higher isolation levels (Read Uncommitted, Read Committed) offer higher concurrency, allowing multiple transactions to execute concurrently. Lower isolation levels (Repeatable Read, Serializable) may introduce more contention and potential for conflicts.
- **Data Integrity:** Higher isolation levels (Repeatable Read, Serializable) provide stronger guarantees for data integrity by preventing non-repeatable reads and phantom reads. Lower isolation levels (Read Uncommitted, Read Committed) may allow dirty reads or inconsistent data.
- **Application Logic:** The specific requirements and logic of your application should guide your choice of isolation level. Consider the criticality of data integrity, the nature of concurrent operations, and the need for predictable query results.

By understanding the trade-offs between concurrency and data integrity, you can select the appropriate isolation level that balances the needs of your application. It's important to test and verify the behaviour of your application under different isolation levels to ensure it meets your requirements for consistency and concurrency.

Resources & References

- [📖 Best practices for creating multi-column indexes in PostgreSQL](#)